

Fitting probabilistic models

Exercise 02

Submitted by:

Hitik Panchal (50230156, s17hpanc)
Rouy Alfahel (3389227, s6roalfa)
Orhan Ugur Aydin (3209243, s6oraydi)
Sebastian Sauerbier (2973879, s6sesaue)
Chenfang Wang, (50292921, s59cwang)

w

Khasha (, S76kalav)
Dengkui Hou (50347144, s77dhou)
Feiyue Cheng (3305865, s6fechen)

Task 2.1: Preliminaries

- Topic: Introduction to the Weibull distribution – implementing and visualizing its PDF and CDF
- Goal: Implement the Weibull PDF (Probability Density Function) and CDF (Cumulative Distribution Function) in NumPy and visualize them as preparation for later modeling tasks
 - i. The code loads and cleans the MySpace dataset, removing leading zeros and creating a time axis to prepare the data for analysis
 - ii. It then visualizes the cleaned data and implements the Weibull PDF and CDF to illustrate and understand the time-based distribution
 - `weibull_pdf(t, α, β)`: computes the density at each time value (probability per unit time) given the parameters α (shape) and β (scale)
 - `weibull_cdf(t, α, β)`: shows how the cumulative probability increases from 0 and approaches 1 as $t \rightarrow \infty$

```
# ---- 4) Weibull PDF/CDF ----
def weibull_pdf(t, alpha, beta):
    """
    Weibull PDF:
        f(t|α,β) = (α/β) * (t/β)^(α-1) * exp( - (t/β)^α ), for t >= 0
    """
    t = np.asarray(t, dtype=float)
    a = float(alpha); b = float(beta)
    if a <= 0 or b <= 0:
        raise ValueError("alpha and beta must be positive.")
    pdf = (a / b) * (t / b) ** (a - 1.0) * np.exp(- (t / b) ** a)
    # Return 0 for negative t
    return np.where(t >= 0.0, pdf, 0.0)

def weibull_cdf(t, alpha, beta):
    """
    Weibull CDF:
        F(t|α,β) = 1 - exp( - (t/β)^α ), for t >= 0
    """
    t = np.asarray(t, dtype=float)
    a = float(alpha); b = float(beta)
    if a <= 0 or b <= 0:
        raise ValueError("alpha and beta must be positive.")
    cdf = 1.0 - np.exp(- (t / b) ** a)
    # Return 0 for negative t
    return np.where(t >= 0.0, cdf, 0.0)
```

Data Visualization and Weibull Modeling

Goal:




Build intuition by visualizing the cleaned MySpace series and the Weibull PDF with the professor's parameter variations (different α and β values)

Result:

- The Weibull family produces rise–peak–decline shapes consistent with the observed MySpace activity; later tasks will fit parameters quantitatively

- Description:

- Middle plot (c vs t): real MySpace counts after removing leading zeros; time axis $t = 1, \dots, n$ by design
- Left plot (Weibull PDFs, vary $\alpha, \beta=2.0$): changing the shape α shifts the peak and skew (larger $\alpha \rightarrow$ peak later, less right-skew)
- Right plot (Weibull PDFs, vary $\beta, \alpha=2.5$): changing the scale β stretches/compresses along the t -axis (larger $\beta \rightarrow$ curve spreads to the right)



Task 2.2 An "Intelligent" Weibull Fit

What was our goal?

- **Goal:** Find the best "knobs" (parameters α, β) for the Weibull distribution.
- **Standard:** To make our model curve $f(t)$ fit the myspace data c as closely as possible.

What method did we use?

- **Maximum Likelihood Estimation (MLE):**
 - A statistical method to find the parameters that make our observed data **most probable**.
 - We didn't just "eyeball" the fit; we found the mathematically "most likely" solution.

Why was this task hard?

- **Challenge: No "Formula" for the Answer**
 - The equations from MLE had **no closed-form solution**.
 - We couldn't just solve for α and β directly like we solve $2x = 10 \implies x = 5$.

Our Solution: "Newton's Method"

- We used **Newton's Method**, a powerful **iterative** algorithm.
- It works like a "smart guesser":
 - i. Start with an **initial guess** ($\alpha = 1, \beta = 1$).
 - ii. Calculate the **Gradient** (direction) and **Hessian** (curvature).
 - iii. Take a "smart step" to a **better** guess.
 - iv. Repeat 20 times until convergence.

The Core Trick: Handling Data Efficiently


- **Problem:** Our data is a histogram (e.g., "100 counts at $t = 150$ "), but the formulas want $N \approx 23,000$ individual points.
- **Naive way:** Create a giant, slow array.
- **Our way:** Use a **weighted sum**.
 - We told the algorithm: "The point $t = 150$ has a **weight of 100**."
 - This lets the calculation run on $n \approx 400$ points, making it **extremely fast**.

Task 2.2 An "Intelligent" Weibull Fit

The Final Result

- **Parameters:** $\hat{\alpha} \approx 2.808$, $\hat{\beta} \approx 215.428$
- **Visualization:**

Generating plot...



- **Conclusion:** The fit is excellent. Newton's method and we successfully found a great model.

Task 2.3: Fitting a Weibull distribution to a histogram (part 2)

- **Goal:** Fit a continuous Weibull PDF $f(t|\alpha, \beta)$ to the discrete Google Trends data (t_j, c_j) .
- Task 2.2 showed that Maximum Likelihood Estimation is computationally demanding.
- Task 2.3 presents an alternative approach:

- We fit a scaled Weibull PDF \tilde{f} by introducing an parameter A .
- **Model:**

$$\tilde{f}(t|A, \alpha, \beta) = A \cdot f(t|\alpha, \beta)$$

- **Tool:** Use the `scipy.optimize.curve_fit` function, which performs Non-linear Least Squares optimization.

The NLS approach is simpler but requires careful setup, especially regarding constraints:

- ① **Model Function Definition:** We needed to define a Python function `scaled_weibull_pdf(t, A, alpha, beta)` to correctly compute the scaled PDF.
- ② **Initial Guesses (p_0):** A starting point for the optimization was required:
 - $A_0 = 1000, \alpha_0 = 2, \beta_0 = 200$.
- ③ **Parameter Bounds:** The Weibull parameters must be positive $(\alpha, \beta \in \mathbb{R}_+)$. We enforced these constraints:
 - $A, \alpha, \beta \in (0, \infty)$

We utilized the `bounds` argument of curve fit:

```
def scaled_weibull_pdf(t, A, alpha, beta):
    return A * weibull_pdf(t, alpha, beta)

t, c = load_myspace_data()
p0 = [1000.0, 2.0, 200.0]
lower = [1e-6, 1e-6, 1e-6]
upper = [np.inf, np.inf, np.inf]

popt, pcov = curve_fit(
    f=scaled_weibull_pdf,
    xdata=t, ydata=c,
    p0=p0,
    bounds=(lower, upper)
)
```

Task 2.3: Fitting a Weibull distribution to a histogram (part 2)

Final Estimated Parameters:

Parameter	Initial Guess	Final Estimate
A (Amplitude)	1000.0	17459.4112
α (Shape)	2.0	2.7310
β (Scale)	200.0	211.0244

Table: Results from the `curve_fit` NLS approach.

Conclusion:

- curve fit is efficient for fitting scaled PDFs.
- The fitted parameters ($\alpha \approx 2.73, \beta \approx 211$) are numerically different from those obtained by other methods (as in Task 2.2).
- Different optimization criteria lead to equally valid models.
- Task 2.3 NLS: Minimizes the Sum of Squared Errors in L2 norm.
- Task 2.2 MLE: Maximizes the Log-Likelihood.




Figure: Raw data fitted with the Weibull curve.

Task 2.4: Fitting a Weibull distribution to a histogram (part 3)

We model the observed count vector $c = (c_1, c_2, \dots, c_n)$ with a multinomial distribution:

$$c \sim \text{Multinomial}(C, p_1, \dots, p_n), C = \sum_{j=1}^n c_j,$$

$$p(c_1, \dots, c_n) = C! \prod_{j=1}^n \frac{p_j^{c_j}}{c_j!},$$

$$p_j(\alpha, \beta) = F(t_j; \alpha, \beta) - F(t_{j-1}; \alpha, \beta),$$

$$F(t | \alpha, \beta) = 1 - \exp\left[-\left(\frac{t}{\beta}\right)^\alpha\right].$$

Objective:

$$\max_{\alpha, \beta} L(\alpha, \beta) = \frac{C!}{F(t_n) - F(t_0)} \prod_{j=1}^n \frac{p_j(\alpha, \beta)^{c_j}}{c_j!}$$

↓ a local quadratic approximation of log-likelihood

$$\min_{\alpha, \beta} \sum_{j=1}^n w_j (c_j - C p_j(\alpha, \beta))^2, \quad w_j = \frac{1}{\sqrt{p_j}}.$$

The numerical optimization problem is approximated by a weighted nonlinear least squares problem, which is solved using Iteratively Reweighted Least Squares (IRLS).

IRLS alternates between:

- ① Update weights w_j given (α, β) .
- ② With the updated weights w_j , compute new parameters (α, β) that minimize the objective (using SCIPY.OPTIMIZE.LEAST_SQUARES).

https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.least_squares.html

The objective function computes the **vector of residuals**.

```
# Residual definition under fixed weights:  
# r_j = w_j * (c_j - C * p_j(theta))  
def residuals(theta):  
    a, b = th  
    p_th = bin_probs_from_cdf(t, a, b)  
    return np.sqrt(w) * (c - C * p_th)  
  
# Positive parameter bounds  
lb = np.array([1e-8, 1e-8])  
ub = np.array([np.inf, np.inf])  
  
res = least_squares(residuals, theta,  
                     bounds=(lb, ub),  
                     xtol=1e-12, ftol=1e-12, gtol=1e-12)  
theta_new = res.x
```

At each outer iteration:

$$\theta^{(k)} = (\alpha^{(k)}, \beta^{(k)}).$$

Residuals for inner least-squares:

$$r_j(\theta) = \sqrt{w_j^{(k)}} (c_j - C p_j(\theta)).$$

Solve:


$$\theta^{(k+1)} = \arg \min_{\theta} \sum_{j=1}^n r_j(\theta)^2.$$

Stop when:

$$\Delta\theta = \|\theta^{(k+1)} - \theta^{(k)}\| < 10^{-8}.$$

Task 2.4: Fitting a Weibull distribution to a histogram (part 3)

Fitted Curve v.s. Data $A = 17370.595583$



Objective:

$$\min_A \sum_{j=1}^n (c_j - Af_j)^2.$$

$$f_j = f(t_j; \hat{\alpha}, \hat{\beta}) = \frac{\hat{\alpha}}{\hat{\beta}} \left(\frac{t_j}{\hat{\beta}} \right)^{\hat{\alpha}-1} \exp \left[- \left(\frac{t_j}{\hat{\beta}} \right)^{\hat{\alpha}} \right].$$

Recall in exercise 1:

$$w_* = \arg \min_w \|X^\top w - y\|^2,$$
$$X^\top = f, \quad w = A, \quad y = c,$$

$$A = \frac{\langle c, f \rangle}{\langle f, f \rangle}$$

Initial parameters:

$$\alpha = 2, \quad \beta = 200.$$

Final estimates:

$$\hat{\alpha} = 2.780828, \quad \hat{\beta} = 212.017505.$$

Task 2.5: Fitting a Weibull model to a discrete distribution (part 1)

Objective: Fit a discrete Weibull model distribution \mathbf{f} to our discrete "myspace" data distribution \mathbf{q} by minimizing the Kullback-Leibler (KL) Divergence.

Methodology:

- **Prepare Data Distribution (\mathbf{q}):** We normalized the raw counts \mathbf{c} and scaled them by **0.98**, as instructed.
 - $\mathbf{q} = (\mathbf{c} / \text{sum}(\mathbf{c})) * 0.98$
 - This vector \mathbf{q} represents our "ground truth" data: a discrete probability distribution (PMF) that sums to **0.98**.
- **Define Model Distribution (\mathbf{f}):**
 - It first calculates the Weibull PDF values for all time steps \mathbf{t} .
 - It then **normalizes** this vector so it sums to 1.0 ($\mathbf{f} = \mathbf{f_vec} / \text{sum}(\mathbf{f_vec})$). This vector \mathbf{f} is our *model's* PMF, which we can compare directly to \mathbf{q} .
- **Define Loss Function:**
 - The goal is to minimize: $D_{KL}(q||f) = \sum q_j \log(\frac{q_j}{f_j})$.
 - We can simplify this. Since the entropy of \mathbf{q} is constant, minimizing the KL divergence is identical to minimizing the **Cross-Entropy**: $\text{Loss} = -\sum q_j \log(f_j)$
- **Optimize:** We used `scipy.optimize.minimize` to find the **alpha** and **beta** that result in the lowest possible loss score.


Loss Function:

```
def kl_loss_function(params: list, t_array: ndarray, q_array: ndarray) -> float:  
    alpha, beta = params  
    model_vector = weibull_pdf(t_array, alpha, beta)  
  
    # Convert model_vector to PMF  
    model_vector_sum = np.sum(model_vector)  
    model_pmf = model_vector / model_vector_sum # normalizing: sums to 1  
  
    model_pmf = np.maximum(model_pmf, 1e-10) # handle numerical instability: avoiding log(0)  
  
    # Loss function = -sum(q_j * log(f_j))  
    loss = -np.sum(q_array * np.log(model_pmf))  
  
    return loss
```

Minimize Function:

```
# Running optimization to minimize KL divergence...  
result = minimize(kl_loss_function, [alpha, beta], args=(t, q), method='L-BFGS-B',  
bounds=parameter_bounds)
```

Task 2.5: Fitting a Weibull model to a discrete distribution (part 1)



Results

- The `scipy.optimize.minimize` function successfully converged.
- Fitted **alpha**: 2.8015
- Fitted **beta**: 215.5443
- Final Loss: 5.5855

Observations

- Minimizing the KL Divergence (via Cross-Entropy) is a powerful method for aligning the *shape* of a theoretical model with an observed data distribution.
- The resulting fit provides a simple, smooth two-parameter model that effectively captures the core growth-and-decline dynamic of the noisy "myspace" trend data.
- This method does not require a final manual scaling step, as we are fitting the *probabilities* directly. However, it relies on the 98% assumption, which is a key modeling decision.

Task 2.6: Coefficient of Variation (Weibull)

Objective:

Fit the Weibull model to a discrete distribution using the method of moments.

The *first-* and *second moment* of the Weibull distribution are the mean and variance:

$$\begin{aligned}\mathbb{E}[T] &= \mu \\ &= \int_0^{\infty} f(t|\alpha, \beta) t dt \\ \mathbb{V}[T] &= \sigma^2 \\ &= \int_0^{\infty} f(t|\alpha, \beta) (t - \mu)^2 dt \\ &= \mathbb{E}[T^2] - \mathbb{E}^2[T]\end{aligned}$$

Task 2.6: Coefficient of Variation (Weibull)

Objective:

Fit the Weibull model to a discrete distribution using the method of moments.

The *first-* and *second moment* of the Weibull distribution are the mean and variance: (*expressed using Gamma function Γ*)

$$\mu = \beta \Gamma(1 + \frac{1}{\alpha})$$

$$\sigma^2 = \beta^2 \left[\Gamma(1 + \frac{2}{\alpha}) - \Gamma^2(1 + \frac{1}{\alpha}) \right]$$

Then the *coefficient of variation* of the Weibull distribution is:

$$CV^2 = \frac{\sigma^2}{\mu^2} = \frac{\Gamma(1 + \frac{2}{\alpha}) - \Gamma^2(1 + \frac{1}{\alpha})}{\Gamma^2(1 + \frac{1}{\alpha})}$$

Task 2.7: Numerically Solving Differential Equations

Goal

Numerically solve the ODE for the Weibull CDF and compare with the analytical CDF and a discrete cumulative from the PDF:

$$\dot{F}(t) = \frac{\alpha}{\beta} \left(\frac{t}{\beta} \right)^{\alpha-1} (1 - F(t)), \quad F(0) = 0$$

Parameters: $\alpha = 2.86$, $\beta = 215.91$, $t \in [0, 450]$

Approach

- Defined Uniform grid t_1, \dots, t_n with $\Delta t = 1$.
- Computed Discrete CDF: $Q_j = \sum_{i=1}^j f(t_i) \Delta t$.
- Derived Analytical CDF:
 $F^a(t) = 1 - \exp[-(t/\beta)^\alpha]$.
- Solved ODE with `solve_ivp` (RK45, `rtol=1e-8, atol=1e-10`).
- Compared $F^n(t)$, $F^a(t)$, and $Q(t)$ to verify accuracy.

Code snippet:

```
# pdf -> q_i -> cumulative Q_j
def weibull_pdf(tt, a, b):
    tt = np.asarray(tt)
    z = np.maximum(tt, 0.0) / b
    return (a / b) * np.power(z, a - 1.0) * np.exp(-np.power(z, a))

q = weibull_pdf(t, alpha, beta) * dt
Q = np.clip(np.cumsum(q), 0.0, 1.0)

# Analytical CDF F^a
def weibull_cdf(tt, a, b):
    tt = np.asarray(tt)
    z = np.maximum(tt, 0.0) / b
    return 1.0 - np.exp(-np.power(z, a))

F_a = weibull_cdf(t, alpha, beta)


# ODE: F'(t) = h(t) * (1 - F), F(0)=0
def hazard(tt, a, b):
    if tt < 0:
        return 0.0
    return (a / b) * ((tt / b) ** (a - 1.0))

def rhs(tt, F):
    return hazard(tt, alpha, beta) * (1.0 - F)

sol = solve_ivp(
    rhs, (t_min, t_max), [0.0],
    t_eval=t, method="RK45", rtol=1e-8, atol=1e-10
)
F_n = sol.y[0]
```

Task 2.7: Numerically Solving Differential Equations

Comparison of Results:



Accuracy:

Comparison	Max $ \Delta $
F_n vs F_a	1.1×10^{-8}
Q vs F_a	2.6×10^{-3}

Observations and Learnings:

- The numerical ODE solution $F^n(t)$ overlaps almost perfectly with the analytical Weibull CDF $F^a(t)$, confirming that our differential formulation is correct.
- The cumulative distribution Q_j , obtained by summing the PDF, also follows the same smooth S-shaped curve, verifying the discrete numerical integration approach.
- Both methods converge to 1 as t increases, which matches the expected CDF behavior.
- Small numerical deviations ($< 10^{-8}$) highlight that the chosen tolerances ($rtol=1e-8$, $atol=1e-10$) were effective and stable.
- The comparison illustrates how ODE solvers like `solve_ivp` can accurately reproduce analytical probability models when parameters are well-defined.
- This task strengthened our understanding of how continuous distributions can be represented as differential equations and verified numerically.

Task 2.8: Neural estimation of Weibull(α , β) for MySpace trend data

Goal: Fit a Weibull distribution to Google Trends data (myspace.csv) using a neural network approach.

Methodology:

- NN with 2 trainable parameter (Pytorch)
- Adam optimizer, 4_000 steps, lr=0.05
- Data q =normalized vounts form myspace.csv
- Loss Function: $L = D_{KL}(f \parallel q) = \sum_j f_j (\log f_j - \log q_j)$

```
class WeibullParamNet(nn.Module):
    def __init__(self, init_alpha=2.0, init_beta=200.0):
        super().__init__()
        self.raw_alpha = nn.Parameter(torch.tensor([math.log(math.exp(init_alpha) - 1.0)], dtype=torch.float32))
        self.raw_beta = nn.Parameter(torch.tensor([math.log(math.exp(init_beta) - 1.0)], dtype=torch.float32))
        self.softplus = nn.Softplus()
    def forward(self):
        alpha = self.softplus(self.raw_alpha) + 1e-6
        beta = self.softplus(self.raw_beta) + 1e-6
        return alpha.squeeze(0), beta.squeeze(0)
```

Task 2.8: Neural estimation of Weibull(α , β) for MySpace trend data

- **Learned parameter:**
 - $\alpha = 2.87$
 - $\beta = 215.25$
- **Observation**
 - Neural network captures sudden rise and decay of the MySpace trend
 - Softplus ensure α , β bigger than zero without explicit constraints
- **Learnings:**
 - Even 2-parameter NN can solve non-linear distribution fitting task
 - Gradient-based optimization works well for statistical parameter estimation

